

Semantic content through encoder then decoder (CLS token, linear probe, 40/10 per class)

